

First, we import a number of libraries that will be used in the analysis

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import os
```

The os library is then used to loop through all the files in the /kaggle/input directory and print the file names. This step is not necessary for the analysis and can be skipped.

Next, we load the heart.csv dataset into a Pandas dataframe called df and display the first 5 rows using the head() method:

```
df = pd.read_csv('heart.csv')
```

```
df.head()
```

We can get some information about the data using the info() method:

```
df.info()
```

We then check for missing values in the data using the isna() method and the sum() method:

```
df.isna().sum()
```

Next, we create a statistical summary of the data using the describe() method and transpose the results using the transpose() method:

```
df.describe().transpose()
```

We can also create a pairplot using the pairplot() method from the seaborn library, setting the hue parameter to HeartDisease:

```
sns.pairplot(df,hue='HeartDisease')
```

We can create a heatmap of correlations between features using the heatmap() method from the seaborn library and setting the figsize parameter to a tuple of (6,4) and the dpi parameter to 150:

```
plt.figure(figsize=(6,4),dpi=150)
```

```
sns.heatmap(df.corr(),cmap='viridis',annot=True)
```

We can find the top 5 correlated features with the target variable by taking the absolute value of the correlations using the abs() method, sorting the values using the sort_values() method, and selecting the last 6 values using the tail() method:

```
np.abs(df.corr()['HeartDisease']).sort_values().tail(6)
```

Next, we split the data into training and test sets using the `train_test_split()` function from the `sklearn.model_selection` library, setting the `test_size` parameter to 0.1 and the `random_state` parameter to 42:

```
from sklearn.model_selection import train_test_split
X = df.drop('HeartDisease',axis=1)
y = df['HeartDisease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

We can check the shape of the training and test sets using the `shape` attribute:

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

continue further i will only type c for continue

Certainly, here is the rest of the code:

We can check the balance of the label data in the training set by using the `value_counts()` method and plotting the results using the `countplot()` function from the `seaborn` library:

```
y_train.value_counts()
sns.countplot(x=y_train)
```

We can calculate the percentage of the minority class in the training set by dividing the number of minority class samples by the total number of samples:

```
372/(454+372)
```

Next, we perform feature engineering by creating a ColumnTransformer object to scale the numeric features and one-hot encode the categorical features. We create lists of the numeric and categorical columns and instantiate a StandardScaler object and a OneHotEncoder object, setting the drop parameter of the OneHotEncoder to 'first' to drop one of the encoded columns:

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
```

```
col_numeric = ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']
col_categoric = ['Sex', 'ChestPainType', 'FastingBS', 'RestingECG', 'ExerciseAngina', 'ST_Slope']
```

```
scaler = StandardScaler()
encoder = OneHotEncoder(drop='first')
```

```
preprocessor = ColumnTransformer([
    ('num', scaler, col_numeric),
    ('cat', encoder, col_categoric)
])
```

We then create a base model of logistic regression by instantiating a LogisticRegression object and setting the max_iter parameter to 1000:

```
from sklearn.linear_model import LogisticRegression
base_model = LogisticRegression(max_iter=1000)
```

We create a pipeline by instantiating a Pipeline object and passing it a list of tuples containing the steps in the pipeline, in this case the preprocessor and the base_model:

```
from sklearn.pipeline import Pipeline
pipe = Pipeline([('preprocessor', preprocessor), ('base_model', base_model)])
```

We perform a grid search using the GridSearchCV function from the sklearn.model_selection library, passing it the pipeline and a dictionary of parameters to test. We set the scoring parameter to 'accuracy' and the cv parameter to 5:

```
from sklearn.model_selection import GridSearchCV
```

```
parameters = ({'base_model__C': [0.001,0.01,0.1,1,10],
               'base_model__penalty': ['l1','l2','elasticnet'],
               'base_model__solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
              })
```

```
grid = GridSearchCV(estimator=pipe,param_grid=parameters,scoring='accuracy',cv=5)
grid.fit(X_train,y_train)
```

We can find the best estimator and best parameters using the `best_estimator_` and `best_params_` attributes of the `grid` object: Certainly, here is the rest of the code:

We can find the best estimator and best parameters using the `best_estimator_` and `best_params_` attributes of the `grid` object:

```
best_model = grid.best_estimator_
best_params = grid.best_params_
print(best_model)
print(best_params)
```

We can calculate the accuracy of the model on the test set using the `score()` method:

```
accuracy = best_model.score(X_test,y_test)
print(accuracy)
```

We can also make predictions using the `predict()` method and calculate the confusion matrix using the `confusion_matrix()` function from the `sklearn.metrics` library:

```
from sklearn.metrics import confusion_matrix
y_pred = best_model.predict(X_test)
confusion_matrix = confusion_matrix(y_test,y_pred)
print(confusion_matrix)
```

Finally, we can plot the confusion matrix using the `heatmap()` function from the seaborn library and setting the `annot` parameter to `True`:

```
plt.figure(figsize=(4,4))  
sns.heatmap(confusion_matrix,annot=True,cmap='Blues',fmt='d')  
plt.ylabel('True label')  
plt.xlabel('Predicted label')
```

This completes the analysis of the heart disease dataset using logistic regression.